

```
# start here without assistance
import seaborn as sns
```

```
def plot_graph_grid(X, y, legend_labels):
```

```
    """
```

```
    plots a grid of scatter plots and KDE plots for the given data X and labels y
```

```
    Parameters:
```

```
        X : the feature data
```

```
        y : the target labels
```

```
        legend_labels : the legend labels corresponding to the target labels
```

```
    Returns:
```

```
        nothing
```

```
    """
```

```
    fig, ax = plt.subplots(4, 4, figsize=(12, 12), sharex='col')
```

```
    for i, feature in enumerate(selected_features):
```

```
        for j, other_feature in enumerate(selected_features):
```

```
            # plot scatter plots for all combinations of features
```

```
            if feature != other_feature:
```

```
                for label in np.unique(y):
```

```
                    # learned about scatterplots from
```

```
https://matplotlib.org/stable/api/\_as\_gen/matplotlib.pyplot.scatter.html
```

```
                    ax[i, j].scatter(X[y == label, selected_features.index(other_feature)], X[y ==
label, selected_features.index(feature)], label=legend_labels[label], edgecolor='white')
```

```
                    # remove tick labels
```

```
                    ax[i, j].tick_params(axis='both', which='both', labelbottom=False, labelleft=True)
```

```
            else:
```

```
                # plot kde plots for the same feature
```

```
                for label in np.unique(y):
```

```
                    # only show x label for the bottom row and y label for the left column
```

```
                    if i == 0:
```

```
                        # learned about seaborn and kde plots from
```

```
https://seaborn.pydata.org/generated/seaborn.kdeplot.html
```

```
                        sns.kdeplot(X[y == label, selected_features.index(feature)], ax=ax[i, j],
label=legend_labels[label], alpha=0.3, fill=True).set(xlabel="", ylabel=feature)
```

```
                    elif j == 3:
```

```
                        sns.kdeplot(X[y == label, selected_features.index(feature)], ax=ax[i, j],
label=legend_labels[label], alpha=0.3, fill=True).set(xlabel=feature, ylabel="")
```

```
                    else:
```

```
                        sns.kdeplot(X[y == label, selected_features.index(feature)], ax=ax[i, j],
label=legend_labels[label], alpha=0.3, fill=True).set(xlabel="", ylabel="")
```

```
                        ax[i, j].tick_params(axis='both', which='both', labelbottom=False, labelleft=False)
```

```
    # set x and y labels for the scatter plots
```

```
    for i in range(len(selected_features)):
```

```
        ax[3, i].set_xlabel(selected_features[i])
```

```

ax[i, 0].set_ylabel(selected_features[i])
ax[i, 0].tick_params(axis='y', labelleft=True, labelbottom=False)
ax[3, i].tick_params(axis='x', labelbottom=True, labelleft=False)

# show tick values on all axes
for ax_row in ax:
    for ax_elem in ax_row:
        ax_elem.tick_params(axis='both', which='both', labelleft=True, labelbottom=True)

# set the legend for all plots
ax[0, 3].legend(loc='upper left', bbox_to_anchor=(1.05, 1))
plt.tight_layout()
plt.show()

df_wine.head()

# run the plotting function
plot_graph_grid(X, y, wine.target_names)

# noise code
mySeed = 12345
np.random.seed(mySeed)
XN=X+np.random.normal(0,0.6,X.shape)

plot_graph_grid(XN, y, wine.target_names)

# helper code

def split_data(X,y):
    """
    splits the given data into training and test sets

    Parameters:
        X : the feature data
        y : the target labels

    Returns:
        X_train : the training feature data
        X_test : the test feature data
        y_train : the training target labels
        y_test : the test target labels
    """
    # set the random seed
    np.random.seed(mySeed)

    samples = len(X)

    indices = np.random.permutation(samples)

```

```
# get the number of training samples (70% of the data)
train_samples = int(np.ceil(samples * 0.7))
```

```
X_train = X[indices[:train_samples]]
y_train = y[indices[:train_samples]]
```

```
X_test = X[indices[train_samples:]]
y_test = y[indices[train_samples:]]
```

```
return X_train, X_test, y_train, y_test
```

```
def distance_between_vectors(x1, x2):
```

```
    """
```

```
    calculates the Euclidean distance between two vectors
```

```
    Parameters:
```

```
        x1 : the first vector
```

```
        x2 : the second vector
```

```
    Returns:
```

```
        the Euclidean distance between the two vectors
```

```
    """
```

```
    distance = np.sqrt(np.sum((x1 - x2)**2))
```

```
    return distance
```

```
def compare_distances(x, X_train, y_train, options):
```

```
    """
```

```
    calculates the Euclidean distance between a vector and all vectors in a dataset and
    returns the shortest distances
```

```
    Parameters:
```

```
        x : the vector to compare distances to
```

```
        X_train : the feature data
```

```
        y_train : the target labels
```

```
        options : the number of shortest distances to return
```

```
    Returns:
```

```
        a list of the shortest distances and their labels
```

```
    """
```

```
    list_of_shortest_distances = []
```

```
    for i in range(len(X_train)):
```

```
        distance = distance_between_vectors(x, X_train[i])
```

```
        if len(list_of_shortest_distances) < options:
```

```
            list_of_shortest_distances.append((distance, y_train[i]))
```

```
        else:
```

```

        list_of_shortest_distances.sort(reverse=True)
        if distance < list_of_shortest_distances[0][0]:
            list_of_shortest_distances[0] = (distance, y_train[i])
    return list_of_shortest_distances

```

```
def mykNN(X,y,X_,options):
```

```
    """
```

Classifies the given data using the k-Nearest Neighbors algorithm.

Parameters:

X : The training feature data.

y : The training target labels.

X_ : The test feature data.

options : The number of neighbors to consider.

Returns:

y_ : The predicted target labels.

```
    """
```

```
    y_ = []
```

```
    for unknown in X_:
```

```
        closest = compare_distances(unknown, X, y, options)
```

```
        for i in range(len(closest)):
```

```
            closest[i] = closest[i][1]
```

```
        #find most common value and if there is a tie, choose the first one
```

```
        #had to freshen up on numpy using numpy docs, reference:
```

```
https://numpy.org/doc/stable/reference/routines.sort.html
```

```
        unique, counts = np.unique(closest, return_counts=True)
```

```
        most_occuring = unique[np.argmax(counts)] #
```

```
        y_.append(most_occuring)
```

```
    return y_
```

```
from sklearn.metrics import accuracy_score
```

```
X_train, X_test, y_train, y_test = split_data(XN, y)
```

```
y_pred = mykNN(X_train, y_train, X_test, 3)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
# confusion matrix, accuracy, precision, recall, etc.
```

```
def confusion_matrix(y_true, y_pred, labels):
```

```
    """
```

calculates the confusion matrix for the given data

Parameters:

y_true : the true target labels

y_pred : the predicted target labels

labels : the labels to use for the confusion matrix

Returns:

cm : the confusion matrix

"""

initialize the confusion matrix

cm = np.zeros((len(labels), len(labels)))

populate the confusion matrix

for i in range(len(y_true)):

cm[y_true[i]][y_pred[i]] += 1

return cm

def draw_confusion_matrix(cm, labels):

"""

draws the confusion matrix via print statements

Parameters:

cm : the confusion matrix

labels : the labels to use for the confusion matrix

Returns:

nothing

"""

print("\tConfusion Matrix")

print("T\P", end="\t")

for label in labels:

print(label, end="\t")

print()

for i in range(len(cm)):

print(labels[i], end="\t")

for j in range(len(cm[i])):

print(cm[i][j], end="\t")

print()

def calc_accuracy(y_true, y_pred):

"""

calculates the accuracy of the given data

Parameters:

y_true : the true target labels

y_pred : the predicted target labels

Returns:

the accuracy of the given data

```

"""
# create the confusion matrix
cm = confusion_matrix(y_true, y_pred, np.unique(y_true))

# calculate the accuracy
correct = cm[0][0] + cm[1][1] + cm[2][2]
total = np.sum(cm)
accuracy = correct / total

return accuracy

def recall(y_true, y_pred, label):
    """
    calculates the recall of the given data for a given label

    Parameters:
        y_true : the true target labels
        y_pred : the predicted target labels
        label : the label to calculate the recall for

    Returns:
        the recall of the given data for the given label
    """
    # create the confusion matrix
    cm = confusion_matrix(y_true, y_pred, np.unique(y_true))

    # calculate the recall
    true_positives = cm[label][label]
    false_negatives = cm[label][0] + cm[label][1] + cm[label][2] - true_positives
    recall = true_positives / (true_positives + false_negatives)

    return recall

def precision(y_true, y_pred, label):
    """
    calculates the precision of the given data

    Parameters:
        y_true : the true target labels
        y_pred : the predicted target labels
        label : the label to calculate the precision for

    Returns:
        the precision of the given data
    """
    # create the confusion matrix
    cm = confusion_matrix(y_true, y_pred, np.unique(y_true))

```

```

# calculate the precision
true_positives = cm[label][label]
false_positives = cm[0][label] + cm[1][label] + cm[2][label] - true_positives
precision = true_positives / (true_positives + false_positives)

return precision

# test evaluation code

print("y_test:", y_test)
print("y_pred:", y_pred)
print()
cm = confusion_matrix(y_test, y_pred, np.unique(y_test))
draw_confusion_matrix(cm, np.unique(y_test))
print()
print("Accuracy:", calc_accuracy(y_test, y_pred))
print("Recall:", recall(y_test, y_pred, 0))
print("Precision:", precision(y_test, y_pred, 0))
print()

#compare to sklearn
from sklearn.neighbors import KNeighborsClassifier
sk_classifier = KNeighborsClassifier(n_neighbors=3)
sk_classifier.fit(X_train, y_train)
y_pred_sk = sk_classifier.predict(X_test)
print("y_pred_sk:", y_pred_sk)
print()
cm = confusion_matrix(y_test, y_pred_sk, np.unique(y_test))
draw_confusion_matrix(cm, np.unique(y_test))
print()
print("SK_Accuracy:", calc_accuracy(y_test, y_pred_sk))
print("SK_Recall:", recall(y_test, y_pred_sk, 0))
print("SK_Precision:", precision(y_test, y_pred_sk, 0))
print()

#Make manhattan distance function
def manhattan_distance(x1, x2):
    """
    calculates the manhattan distance between two points

    Parameters:
        x1 : the first point
        x2 : the second point

    Returns:
        the manhattan distance between the two points
    """
    distance = np.sum(np.abs(x1 - x2))

```

```
return distance
```

```
#redo compare distance function to use manhattan distance
```

```
def compare_distances_manhattan(x, X_train, y_train, options):
```

```
    """
```

```
    calculates the Manhattan distance between a vector and all vectors in a dataset and
    returns the shortest distances
```

```
Parameters:
```

```
    x : the vector to compare distances to
```

```
    X_train : the feature data
```

```
    y_train : the target labels
```

```
    options : the number of shortest distances to return
```

```
Returns:
```

```
    a list of the shortest distances and their labels
```

```
    """
```

```
list_of_shortest_distances = []
```

```
for i in range(len(X_train)):
```

```
    distance = manhattan_distance(x, X_train[i])
```

```
    if len(list_of_shortest_distances) < options:
```

```
        list_of_shortest_distances.append((distance, y_train[i]))
```

```
    else:
```

```
        list_of_shortest_distances.sort(reverse=True)
```

```
        if distance < list_of_shortest_distances[0][0]:
```

```
            list_of_shortest_distances[0] = (distance, y_train[i])
```

```
return list_of_shortest_distances
```

```
#redo mykNN to use manhattan distance
```

```
def mykNNManhattan(X, y, X_, options):
```

```
    """
```

```
    Classifies the given data using the k-Nearest Neighbors algorithm.
```

```
Parameters:
```

```
    X : The training feature data.
```

```
    y : The training target labels.
```

```
    X_ : The test feature data.
```

```
    options : The number of neighbors to consider.
```

```
Returns:
```

```
    y_ : The predicted target labels.
```

```
    """
```

```
y_ = []
```



```

for unknown in X_:
    closest = compare_distances_manhattan(unknown, X, y, options)
    for i in range(len(closest)):
        closest[i] = closest[i][1]
    #find most common value and if there is a tie, choose the first one
    unique, counts = np.unique(closest, return_counts=True)
    most_occurring = unique[np.argmax(counts)]
    y_.append(most_occurring)
return y_

```

myNestedCrossVal code

```

def myNestedCrossVal(X, y, k, neighbor_amount, methods, seed):
    """
    performs nested cross validation on the given data and prints the results

```

Parameters:

X : the feature data
 y : the target labels
 k : the number of folds
 neighbor_amount : the number of neighbors
 methods : the distance methods to use
 seed : the seed to use for the random number generator

Returns:

all_accuracies : a list of the accuracies for each fold

"""

```

np.random.seed(seed)
samples = len(X)
indices = np.random.permutation(samples)

```

Shuffle the data

```

X = X[indices]
y = y[indices]

```

Split the data into k folds

```

X_folds = np.array_split(X, k)
y_folds = np.array_split(y, k)

```

```

all_accuracies = []
all_cm = []

```

```

for i in range(k):

```

```

    fold_accuracies = []

```

Create the training and testing data

```

X_train = np.concatenate(X_folds[:i] + X_folds[i+1:])

```

```

y_train = np.concatenate(y_folds[:i] + y_folds[i+1:])
X_test = X_folds[i]
y_test = y_folds[i]

best_method = ""
best_neighbors = 0
best_accuracy = 0

for method in methods:
    # Split the training data into k folds
    X_train_folds = np.array_split(X_train, k)
    y_train_folds = np.array_split(y_train, k)

    # Create the training and validation data
    X_train_inner = np.concatenate(X_train_folds[:0] + X_train_folds[1:])
    y_train_inner = np.concatenate(y_train_folds[:0] + y_train_folds[1:])
    X_val_inner = X_train_folds[0]
    y_val_inner = y_train_folds[0]

    if method == "euclidean":
        for neighbors in neighbor_amount:
            # Use my old code to create the model since it uses Euclidean distance
            y_pred_euc = mykNN(X_train_inner, y_train_inner, X_val_inner, neighbors)
            # Calculate the accuracy
            accuracy = calc_accuracy(y_val_inner, y_pred_euc)
            # Check if this is the best accuracy in the fold
            if accuracy > best_accuracy:
                best_accuracy = accuracy
                best_neighbors = neighbors
                best_method = method
            # Add the accuracy to the list
            fold_accuracies.append(accuracy)
            print(f"Fold {i+1} - {method} - Neighbors={neighbors}- Accuracy:", accuracy)
    elif method == "manhattan":
        for neighbors in neighbor_amount:
            # Use my Manhattan code to create the model
            y_pred_man = mykNNManhattan(X_train, y_train, X_test, neighbors)
            # Calculate the accuracy
            accuracy = calc_accuracy(y_test, y_pred_man)
            # Check if this is the best accuracy in the fold
            if accuracy > best_accuracy:
                best_accuracy = accuracy
                best_neighbors = neighbors
                best_method = method
            # Add the accuracy to the list
            fold_accuracies.append(accuracy)
            print(f"Fold {i+1} - {method} - Neighbors={neighbors} - Accuracy:", accuracy)

print()

```

```

print(f"Best Accuracy in Fold {i+1}:", best_accuracy)
print(f"Best Neighbors in Fold {i+1}:", best_neighbors)
print(f"Best Method in Fold {i+1}:", best_method)
print()

if best_method == "euclidean":
    y_pred = mykNN(X_train, y_train, X_test, best_neighbors)
    accuracy = calc_accuracy(y_test, y_pred)
    all_accuracies.append(accuracy)
elif best_method == "manhattan":
    y_pred = mykNNManhattan(X_train, y_train, X_test, best_neighbors)
    accuracy = calc_accuracy(y_test, y_pred)
    all_accuracies.append(accuracy)

# Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred, np.unique(y))
print(end="\t")
print(f"Fold {i+1}")
draw_confusion_matrix(cm, np.unique(y))
all_cm.append(cm)
print()
print()

print("OVERALL SUMMARY CONFUSION MATRIX")

#combine all folds into one confusion matrix
overall_cm = np.zeros((len(np.unique(y)), len(np.unique(y))))
for i in range(len(all_accuracies)):
    overall_cm += all_cm[i]
overall_cm = overall_cm.astype(int)
draw_confusion_matrix(overall_cm, np.unique(y))
print()

return all_accuracies

# evaluate clean data code
myNestedCrossVal(X, y, 5, list(range(1,11)), ["euclidean", "manhattan"], mySeed)

# evaluate noisy data code
myNestedCrossVal(XN, y, 5, list(range(1,11)), ["euclidean", "manhattan"], mySeed)

print('CLEAN')
# clean data summary results
myNestedCrossVal(X, y, 5, list(range(1,11)), ["euclidean", "manhattan"], mySeed)

print('NOISY')
# noisy data summary results

```

```
myNestedCrossVal(XN, y, 5, list(range(1,11)), ["euclidean", "manhattan"], mySeed)
```

```
#end here without assistance
```